

Learning to Negotiate: Fine-Tuning and Reinforcement Learning for LLM-Based Bargaining Agents

Dhruv Bansal (dbansa11@asu.edu) Bradley Breisch (bbreisc1@asu.edu)
Luan Nguyen (ltnguy58@asu.edu)

Spring 2026

¹

References

- [1] M. Lewis, D. Yarats, Y. N. Dauphin, D. Parikh, and D. Batra, “Deal or no deal? End-to-end learning for negotiation dialogues,” *arXiv preprint arXiv:1706.05125*, 2017.
- [2] H. He, D. Chen, A. Balakrishnan, and P. Liang, “Decoupling strategy and generation in negotiation dialogues,” in *Proceedings of the 2018 conference on empirical methods in natural language processing*, Association for Computational Linguistics, 2018, pp. 2333–2343.
- [3] H. Touvron *et al.*, “Measuring bargaining abilities of LLMs: An analysis of language model negotiation skills,” *arXiv preprint*, 2024.
- [4] C. R. Team, “EvoEmo: Evolved emotional policies for LLM agents in multi-turn negotiation,” 2025.
- [5] K. Chawla, J. Ramirez, R. Clever, G. Lucas, J. May, and J. Gratch, “CaSiNo: A corpus of campsite negotiation dialogues for automatic negotiation systems,” in *Proceedings of the 2021 conference of the north american chapter of the association for computational linguistics: Human language technologies*, 2021, pp. 3167–3185.
- [6] F. Bianchi *et al.*, “NegotiationArena: A platform for competitive LLM negotiation,” in *Proceedings of the 41st international conference on machine learning*, in ICML’24. 2024.
- [7] A. R. Team, “AgenticPay: A comprehensive framework for language-mediated negotiation,” *arXiv preprint arXiv:2602.06008*, 2026.
- [8] Z. Shao *et al.*, “DeepSeekMath: Pushing the limits of mathematical reasoning in open language models,” *arXiv preprint*, 2024.
- [9] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms,” in *arXiv preprint arXiv:1707.06347*, 2017.
- [10] A. Yang *et al.*, “Qwen technical report,” *arXiv preprint arXiv:2309.16609*, 2024.
- [11] J. Wei *et al.*, “Chain-of-thought prompting elicits reasoning in large language models,” *arXiv preprint arXiv:2201.11903*, 2023.
- [12] Hugging Face, “TRL: Transformer reinforcement learning.” <https://github.com/huggingface/trl>, 2023.
- [13] E. J. Hu, Y. Shen, P. Wallis, *et al.*, “LoRA: Low-rank adaptation of large language models,” *arXiv preprint arXiv:2106.09685*, 2021.
- [14] D. DeVault, J. Mell, and J. Gratch, “Toward natural turn-taking in a virtual human negotiation agent,” in *AAAI spring symposium on turn-taking and coordination in human-machine interaction*, Stanford, CA: AAAI Press, 2015.
- [15] D. Traum, S. C. Marsella, J. Gratch, J. Lee, and A. Hartholt, “Multi-party, multi-issue, multi-strategy negotiation for multi-modal virtual agents,” in *Proceedings of the 8th international conference on intelligent virtual agents*, in IVA’08. Berlin, Heidelberg: Springer-Verlag, 2008, pp. 117–130.

¹Code available at: <https://github.com/dhruvb26/CSE485-Capstone>